

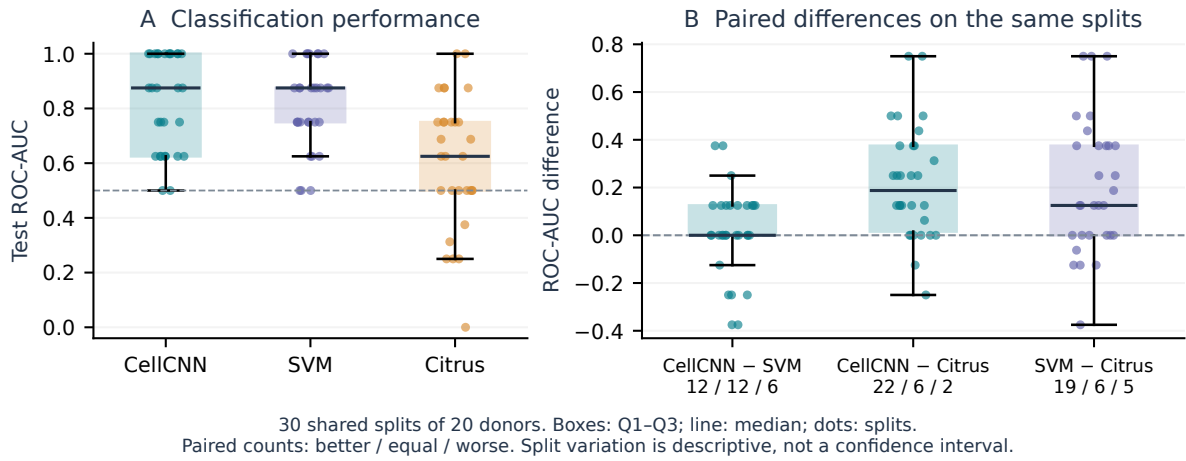


Figure 1: Donor-level classification across 30 shared splits of the same 20 donors, with six test donors per split. Left: test ROC-AUC; right: differences between methods evaluated on the same split. Boxes indicate the interquartile range, horizontal lines the median, and dots individual splits. Counts below paired comparisons denote better/equal/worse splits for the first method. Repeated splits describe sensitivity to donor partitioning; they are not independent cohorts or confidence intervals.

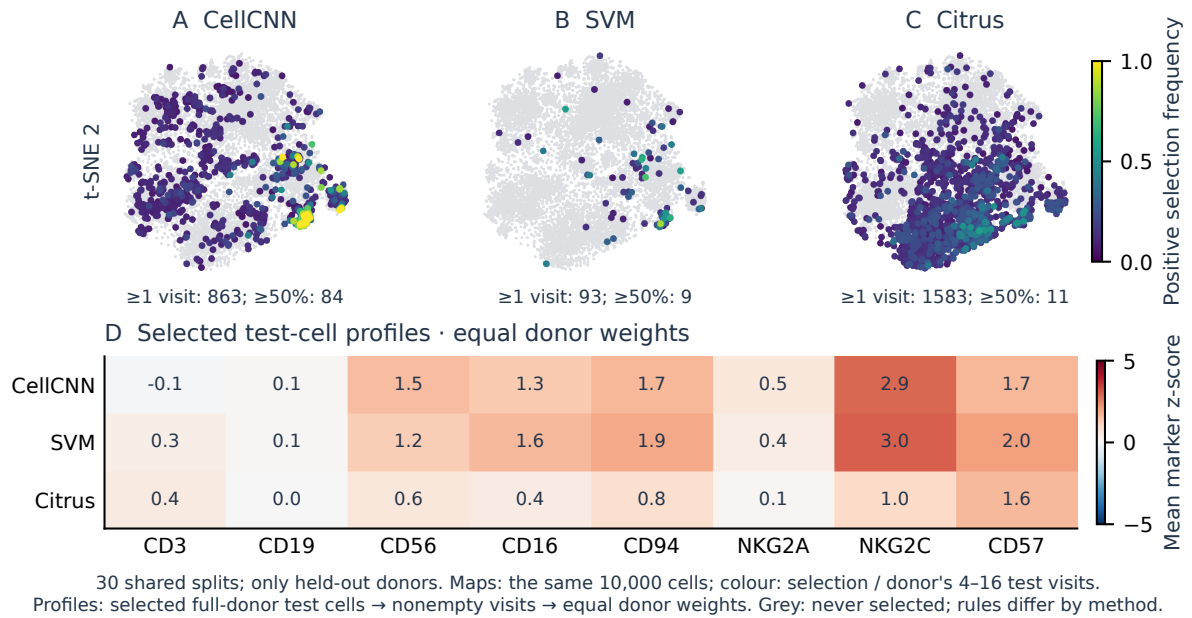


Figure 2: Cell subsets selected in the CMV-positive direction across 30 shared splits. A–C: selection frequency over each donor's 4–16 outer test appearances, shown on the same `gated_alive` t-SNE reference (10,000 cells, 500 per donor, perplexity 30). Labels count cells selected at least once or in at least half of their test appearances. Grey cells were never selected. D: marker profiles of selected full-donor test cells, averaged over nonempty visits within donors and then equally across contributing donors, on a shared reference z-scale. All three methods use test donors only; their selection rules remain method-specific.