

Benchmark of comparative single-cell analysis approaches

Group-Project SSBI 2026

Sebastian Fay

Marie Schygulla

Gregor Habitzreither

September 2026

Introduction

Analysing single-cell data brings some particular difficulties compared to bulk measurements: cells from the same donor are not independent replicates of that donor’s biology, and only a subset of a donor’s cells may show a phenotype of interest, so donor-level labels cannot simply be treated as cell-level labels. In this project we compare and try to improve CellCnn [1], a supervised representation-learning method that uses donor-level labels to detect disease-associated cell subsets.

Our data come from Horowitz et al. [2] and cover 20 donors (9 CMV+, 11 CMV−), including 4 identifiable monozygotic twin pairs. Each donor’s history of cytomegalovirus (CMV) infection was recorded, since CMV infection is known to leave behind a population of long-lived, memory-like ("adaptive") natural killer (NK) cells. This CMV label was assigned donor-wise, even though only a small subpopulation of that donor’s NK cells actually carries the phenotypic signature linked to the infection. From each sample, 37 protein markers were measured by CyTOF, a technique that reads out metal-isotope-tagged antibodies via time-of-flight mass spectrometry. Before subsampling, NK cells made up between 2.4 % and 16.2 % of alive events per donor.

To avoid a bias arising from donors contributing very different numbers of cells, we randomly selected 2,000 cells from each donor to obtain a balanced subsample. To check representativeness, we compared the fraction of NKG2C⁺/CD57⁺ cells (a marker combination associated with the adaptive, CMV-associated NK phenotype) in the balanced subsample against the fraction in the complete cell population of each donor (Supp. Fig. 2): the subsample is representative for this specific two-marker phenotype, though we did not check other marker combinations or rarer populations.

Before dimensionality reduction, we applied an arcsinh transformation with a cofactor of 5, the standard default for CyTOF data [3, 4], which we validated visually via clean bimodal separation of known positive/negative markers (CD3, CD56, CD16, CD57) relative to smaller and larger candidate cofactors. We then z-scored each marker to give every marker equal weight in the distance-based methods below.

Dimensionality reduction

We started with principal component analysis (PCA). As shown in Fig. 1, cumulative explained variance grows roughly like a square-root function of the number of components, so the first two PCs together explain only 20.29 % of the total variance (Supp. Fig. 3). We use the first 29 PCs (91.5 % of the variance) as input to the non-linear methods below to reduce dimensionality before the more expensive computations while retaining most of the variance.

Next, we swept the perplexity parameter of t-SNE [5] (via scikit-learn [6]) over the values 5, 30, 50, 60, 70 and 100, computing trustworthiness and kNN-preservation on a fixed random subsample of 5,000 of our 40,000 balanced cells. trustworthiness penalizes embeddings that place points close together that were not close in the high-dimensional input, weighted by how far out of the true neighbourhood a false neighbour originally ranked [7]; kNN-preservation is stricter, measuring the exact overlap of each point’s k nearest neighbours between the two representations, with no partial credit for a near-miss. trustworthiness scores for perplexities




Figure 1: Cumulative explained variance (blue) across the 36 displayed principal components.

50–100 spanned about 0.0020, so we chose the maximum, 60 (trustworthiness= 0.926). Visually, perplexity 5 gives a largely undifferentiated cloud, while higher perplexities show distinct clusters emerging (Supp. Fig. 4).

For UMAP [8] we grid-searched learning rate (0.1, 0.5, 1.0, 5.0), minimum distance (0.0, 0.1, 0.5, 1) and number of neighbours (5, 15, 50, 100). The highest trustworthiness was obtained at learning rate 0.1, 5 neighbours and minimum distance 0.0 (Table 3, full grid in Supp. Fig. 5). Visibly lower minimum distance packs points more tightly and the number of neighbours changes cluster shape but has a smaller effect than minimum distance.

We compare these three methods using trustworthiness, kNN-preservation and the Pearson correlation (r_d) of all pairwise distances within a random subsample of 2,000 cells. While the first two measures capture *local* structure (is each point’s immediate neighbourhood preserved), r_d captures *global* structure (is the overall arrangement of near and far points preserved, independent of exact neighbour identity). This distinction shows in the results (Table 1): trustworthiness and kNN-preservation agree on t-SNE as the best method, but r_d favours PCA instead, consistent with PCA being a linear projection explicitly built to preserve variance/distance.

Embedding	$T(15) \uparrow$	$R(15) \uparrow$	$r_d \uparrow$
PCA	0.7258	0.0072	0.6204
t-SNE	0.9612	0.2232	0.4868
UMAP	0.9009	0.1126	0.3954

Table 1: PCA, t-SNE and UMAP: local structure via trustworthiness (T) and kNN-preservation (R), global structure via Pearson correlation (r_d) of pairwise distances. T, R use $k = 15$. Green: best per column; red: worst.

Pair	$R(15) \uparrow$	$M^2 \downarrow$
PCA vs t-SNE	0.00585	0.45189
PCA vs UMAP	0.00410	0.73390
t-SNE vs UMAP	0.12532	0.72369

Table 2: Pairwise kNN-preservation (R , local) and Procrustes disparity (M^2 , aligned global shape).

Finally, we compare each pair of methods using kNN-preservation (R) and Procrustes disparity (M^2) (Table 2). Procrustes disparity aligns two maps by translation/rotation/reflection/scaling, then measures the remaining mismatch after this optimal alignment; a lower score means a closer overall shape once cosmetic differences in orientation and scale have been removed [9]. By this measure, PCA and t-SNE are the closest match. By contrast, for local-neighbour agreement (R), t-SNE and UMAP are the more similar pair. The absolute R values look small, but chance level for $n = 40,000$, $k = 15$ is only $k/n \approx 0.0004$, so even the smallest value is well above chance.

References

- [1] Eirini Arvaniti and Manfred Claassen. Sensitive detection of rare disease-associated cell subsets via representation learning. *Nature Communications*, 8:14825, 2017. doi: 10.1038/ncomms14825. URL <https://doi.org/10.1038/ncomms14825>.
- [2] Amir Horowitz, Dara M. Strauss-Albee, Michael Leipold, Jessica Kubo, Neda Nemat-Gorgani, Ozge C. Dogan, Cornelia L. Dekker, Sally Mackey, Holden Maecker, Gary E. Swan, Mark M. Davis, Paul J. Norman, Lisbeth A. Guethlein, Manisha Desai, Peter Parham, and Catherine A. Blish. Genetic and environmental determinants of human NK cell diversity revealed by mass cytometry. *Science Translational Medicine*, 5(208):208ra145, 2013. doi: 10.1126/scitranslmed.3006702. URL <https://doi.org/10.1126/scitranslmed.3006702>.

- [3] Sean C. Bendall, Erin F. Simonds, Peng Qiu, El-ad David Amir, Peter O. Krutzik, Rachel Finck, Robert V. Bruggner, Rachel Melamed, Angelica Trejo, Olga I. Ornatsky, Robert S. Balderas, Sylvia K. Plevritis, Karen Sachs, Dana Pe'er, Scott D. Tanner, and Garry P. Nolan. Single-cell mass cytometry of differential immune and drug responses across a human hematopoietic continuum. *Science*, 332(6030):687–696, 2011. doi: 10.1126/science.1198704. URL <https://doi.org/10.1126/science.1198704>.
- [4] Malgorzata Nowicka, Carsten Krieg, Helena L. Crowell, Lukas M. Weber, Felix J. Hartmann, Silvia Guglietta, Burkhard Becher, Mitchell P. Levesque, and Mark D. Robinson. CyTOF workflow: differential discovery in high-throughput high-dimensional cytometry datasets. *F1000Research*, 6:748, 2017. doi: 10.12688/f1000research.11622.3. URL <https://doi.org/10.12688/f1000research.11622.3>. Discusses and recommends the arcsinh cofactor of 5 for CyTOF data.
- [5] Laurens van der Maaten and Geoffrey Hinton. Visualizing data using t-SNE. *Journal of Machine Learning Research*, 9(86):2579–2605, 2008. URL <https://www.jmlr.org/papers/v9/vandemaaten08a.html>.
- [6] Fabian Pedregosa, Gaël Varoquaux, Alexandre Gramfort, Vincent Michel, Bertrand Thirion, Olivier Grisel, Mathieu Blondel, Peter Prettenhofer, Ron Weiss, Vincent Dubourg, Jake Vanderplas, Alexandre Passos, David Cournapeau, Matthieu Brucher, Matthieu Perrot, and Édouard Duchesnay. Scikit-learn: Machine learning in Python. *Journal of Machine Learning Research*, 12:2825–2830, 2011. URL <https://www.jmlr.org/papers/v12/pedregosa11a.html>. PCA/t-SNE implementations used here: <https://scikit-learn.org/stable/modules/generated/sklearn.decomposition.PCA.html>, <https://scikit-learn.org/stable/modules/generated/sklearn.manifold.TSNE.html>.
- [7] Jarkko Venna and Samuel Kaski. Neighborhood preservation in nonlinear projection methods: An experimental study. In Georg Dorffner, Horst Bischof, and Kurt Hornik, editors, *Artificial Neural Networks — ICANN 2001*, volume 2130 of *Lecture Notes in Computer Science*, pages 485–491, Berlin, Heidelberg, 2001. Springer. doi: 10.1007/3-540-44668-0_68. URL https://doi.org/10.1007/3-540-44668-0_68. scikit-learn implementation used here: <https://scikit-learn.org/stable/modules/generated/sklearn.manifold.trustworthiness.html>.
- [8] Leland McInnes, John Healy, and James Melville. UMAP: Uniform manifold approximation and projection for dimension reduction. *arXiv preprint arXiv:1802.03426*, 2018. URL <https://arxiv.org/abs/1802.03426>. Reference implementation: <https://umap-learn.readthedocs.io/>.
- [9] Pauli Virtanen, Ralf Gommers, Travis E. Oliphant, et al. SciPy 1.0: Fundamental algorithms for scientific computing in Python, 2020. URL <https://doi.org/10.1038/s41592-019-0686-2>. `scipy.spatial.procrustes` documentation: <https://docs.scipy.org/doc/scipy/reference/generated/scipy.spatial.procrustes.html>.

Supplementary Figures and Tables




Figure 2: NKG2C⁺/CD57⁺ fraction, 2,000-cell subsample per donor vs. full population.




Figure 3: First two PCs, 40,000-cell sample, coloured by CMV status (CMV+: red, CMV-: blue).




Figure 4: t-SNE embeddings, perplexities 5, 30, 50, 60, 70, 100; trustworthiness/kNN-preservation above each panel.



Figure 5: UMAP grid: effect of $n_{\text{neighbors}}$ and min_dist at learning rate 0.1.

learning_rate	$n_{\text{neighbors}}$	min_dist	$T(15) \uparrow$	$R(15) \uparrow$
0.1	5	0.0	0.895908	0.230373
0.5	5	0.0	0.893360	0.226933
1.0	5	0.0	0.890139	0.224227
5.0	5	0.0	0.878033	0.203747
0.1	50	0.1	0.864503	0.149040
0.5	50	0.1	0.862615	0.143093
1.0	50	0.1	0.856671	0.135920
5.0	50	0.1	0.840111	0.111320

Table 3: Effect of UMAP learning rate on trustworthiness/kNN-preservation for two fixed ($n_{\text{neighbors}}$, min_dist) settings.